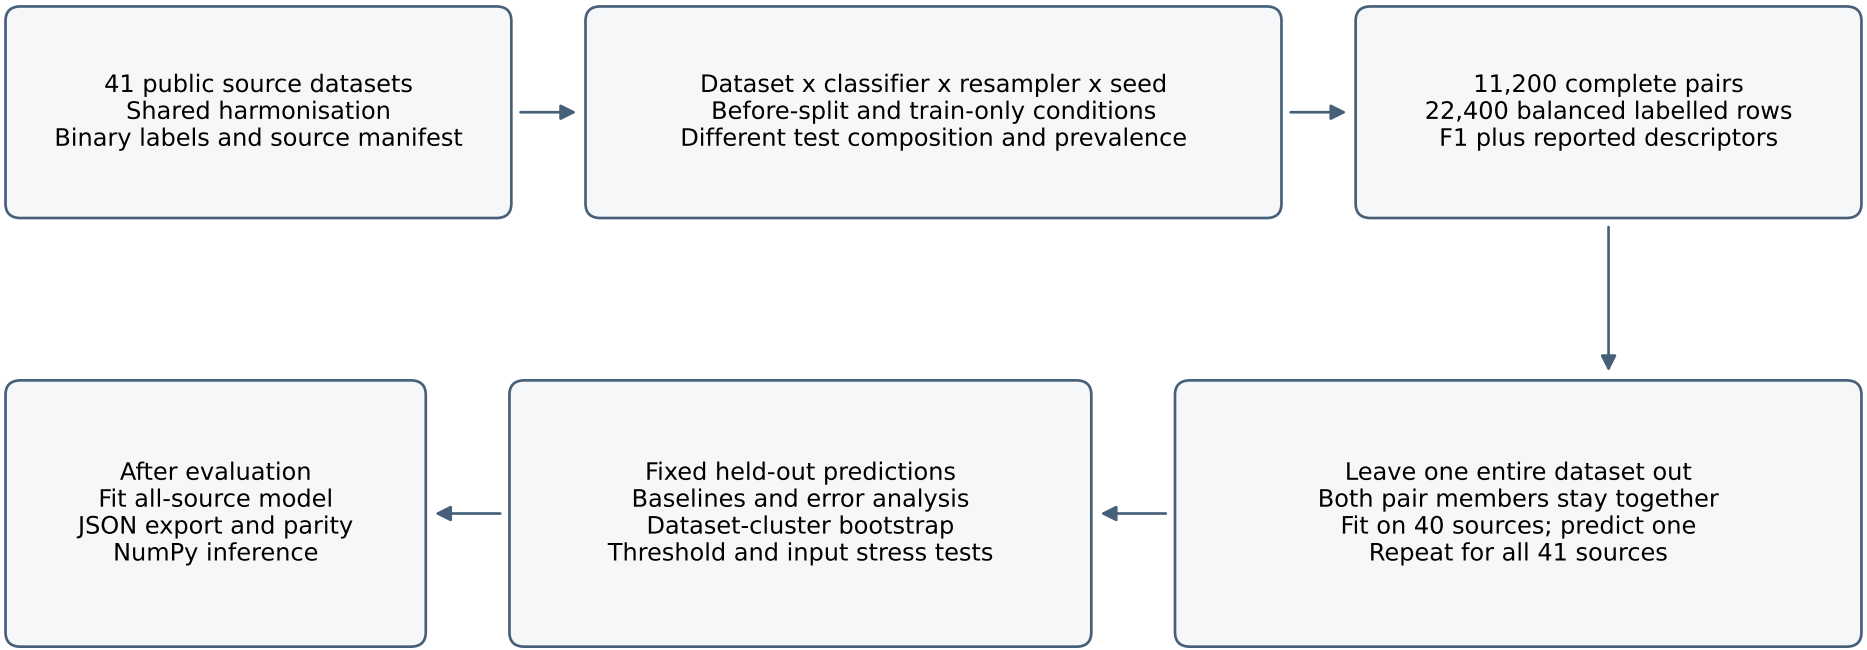
41 public source datasets
Shared harmonisation
Binary labels and source manifest
Dataset x classifier x resampler x seed
Before-split and train-only conditions
Different test composition and prevalence
11,200 complete pairs
22,400 balanced labelled rows
F1 plus reported descriptors
Leave one entire dataset out
Both pair members stay together
Fit on 40 sources; predict one
Repeat for all 41 sources
Fixed held-out predictions
Baselines and error analysis
Dataset-cluster bootstrap
Threshold and input stress tests
After evaluation
Fit all-source model
JSON export and parity
NumPy inference